

Microprocessor Applications

(Term Project 2)



TP (2)

2025. 6. 16.

Students

202110410 조민우

202110365 김상만

- **64pt FFT Accelerator with Multiple Cores**

- **Result of TP_2**
- **Code Review**
- **How to Debug**
- **Ablation Study**
- **Roofline Model**
- **How to improve**
- **Summary**

Result of TP_2

■ Accuracy & Performance

Output count has reached the total number of samples.
Validating bit-accurate output...

- Total mismatches: 0 / 192

Checking SRAM occupancy...

- SRAM 2: occupancy = 32

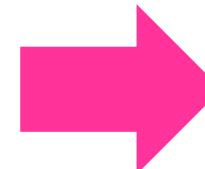
- SRAM 3: occupancy = 32

Test passed. Ending simulation.

\$finish called from file "[../../tb/tb_mem_IO.v](#)", line 132.

\$finish at simulation time [6634000](#)

Simulation complete, time is [6634000 ps](#).



• Accuracy

- Total mismatches : 0/192

• Performance

- Simulation Time : 6,634,000[ps]

• Memory Capacity

- Used Memory: SRAM 0: 32(128B)

- Used Memory: SRAM 1: 32(128B)

- **Implement POV**

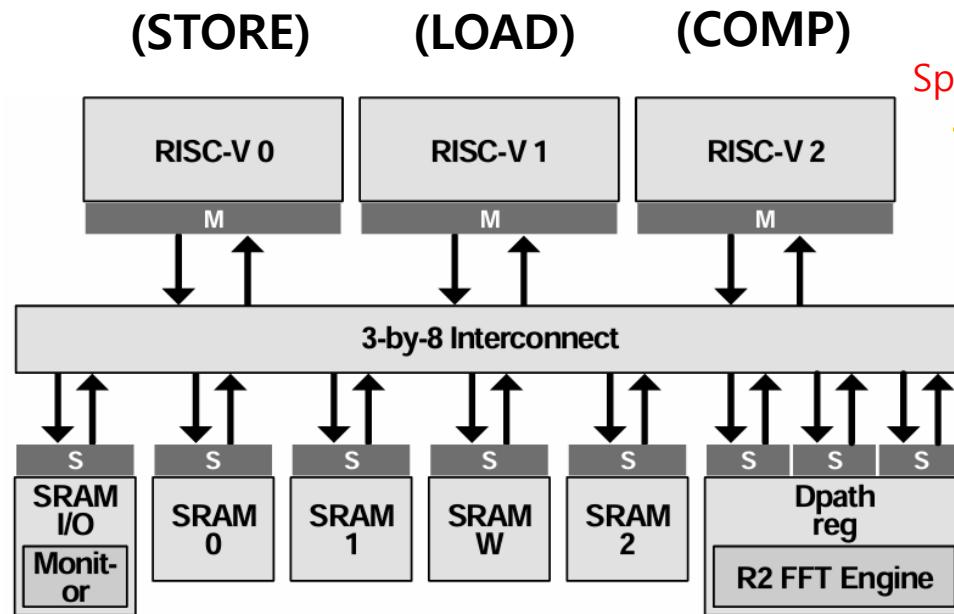
- Pipeline
- Effective Memory Capacity
- Synchronizing with dummy

- **Code Review**

Code Review (1/7)

■ Implement POV

- Pipelining with 3 cores (HOW?)
 - Distribute the roles among the cores



Specify Cores roles

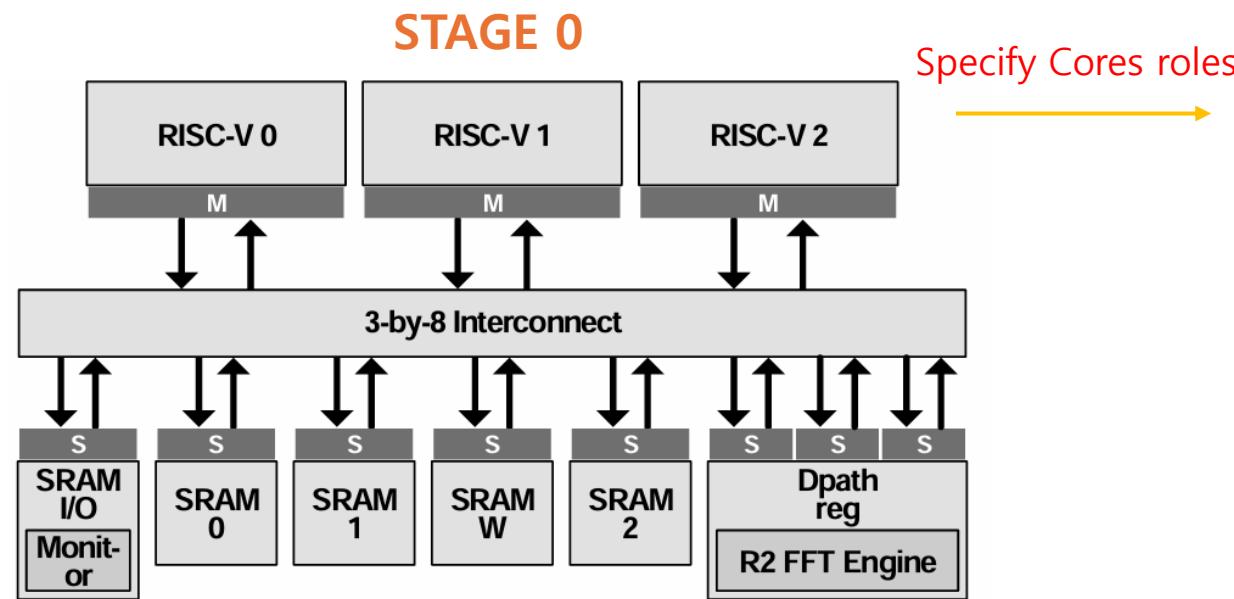
| | A | B | C | D | E | F | G | H | I | J | K | L | M | N | O | P | Q | R | S | T | U | V | W | X | Y | Z | AA | AB | AC | AD | AE | AF | AG | AH |
|---------------|-----------------------------|----------|-----------|-----------|-----------|-----------|-----------|-----------|-----------|-----------|-----------|-----------|-----------|-----------|-----------|-----------|-----------|-----------|-----------|-----------|-----------|-----------|-----------|-----------|-----------|-----------|-----------|-----------|-----------|-----------|-----------|-----------|-----------|-----------|
| stage0 | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | |
| 1 | time-> | 1 | 2 | 3 | 4 | 5 | 6 | 7 | 8 | 9 | 10 | 11 | 12 | 13 | 14 | 15 | 16 | 17 | 18 | 19 | 20 | 21 | 22 | 23 | 24 | 25 | 26 | 27 | 28 | 29 | 30 | 31 | 32 | 33 |
| 2 | risc-v 1 | lw d0 | lw d1 | sw d0 | sw d1 | lw d0 | lw d1 | sw d0 | sw d1 | lw d0 | lw d1 | sw d0 | sw d1 | lw d0 | lw d1 | sw d0 | sw d1 | lw d0 | lw d1 | sw d0 | sw d1 | lw d0 | lw d1 | sw d0 | sw d1 | lw d0 | lw d1 | sw d0 | sw d1 | lw d0 | lw d1 | sw d0 | sw d1 | |
| 3 | risc-v 2 | | | | | lw d0 | lw d1 | sw d0 | sw d1 | |
| 4 | risc-v 0 | lw w(| | | | | | | | sw w(t | sw w(o | sw w(o | sw w(t | sw w(o |
| 5 | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | |
| 6 | stage1,2,3,4,5 | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | |
| 7 | time-> | 1 | 2 | 3 | 4 | 5 | 6 | 7 | 8 | 9 | 10 | 11 | 12 | 13 | 14 | 15 | 16 | 17 | 18 | 19 | 20 | 21 | 22 | 23 | 24 | | | | | | | | | |
| 8 | risc-v 1 | | | | | lw o_ | sw o_ | | | | | | | | | | | |
| 9 | risc-v 2 | lw d0 | lw d1 | sw d0 | sw d1 | lw d0 | lw d1 | sw d0 | sw d1 | lw d0 | lw d1 | sw d0 | sw d1 | lw d0 | lw d1 | sw d0 | sw d1 | lw d0 | lw d1 | sw d0 | sw d1 | lw d0 | lw d1 | sw d0 | sw d1 | lw d0 | lw d1 | sw d0 | sw d1 | lw d0 | lw d1 | sw d0 | sw d1 | |
| 10 | risc-v 0 | lw w(| sw w(t | lw w(o | sw w(o | sw w(t | sw w(o | |
| 11 | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | |
| 12 | stage 5 Output store | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | |
| 13 | time-> | 1 | 2 | 3 | 4 | 5 | 6 | 7 | 8 | 9 | 10 | 11 | 12 | 13 | 14 | 15 | 16 | 17 | 18 | 19 | 20 | 21 | 22 | 23 | 24 | | | | | | | | | |
| 14 | risc-v 1 | | lw o_ | sw o_ | | |
| 15 | risc-v 2 | lw o_ | sw o_ | lw o_ | sw o_ | lw o_ | sw o_ | lw o_ | sw o_ | lw o_ | sw o_ | lw o_ | sw o_ | lw o_ | sw o_ | lw o_ | sw o_ | lw o_ | sw o_ | lw o_ | sw o_ | lw o_ | sw o_ | lw o_ | sw o_ | lw o_ | sw o_ | lw o_ | sw o_ | lw o_ | sw o_ | | | |
| 16 | risc-v 0 | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | |
| 17 | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | |
| 18 | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | | |

Code Review (2/7)

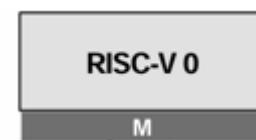
Implement POV-cont'd

- Pipelining with 3 cores (HOW?)

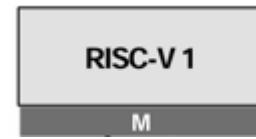
- Distribute the roles among the cores



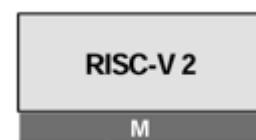
Specify Cores roles



```
# Load twiddle factors for all stages (32 values)
# Butterfly 0-31: Load W[0] for stage 0 (all W[0])
```



```
# Butterfly 0
lw t0, 0(s0)      # d0(SRAM_I/O->t0)
lw t1, 128(s0)    # d1(SRAM_I/O->t1)
sw t0, 0(s1)      # d0(t0->SRAM_0)
sw t1, 0(s2)      # d1(t1->SRAM_1)
addi x0, x0, 1    # dummy
```



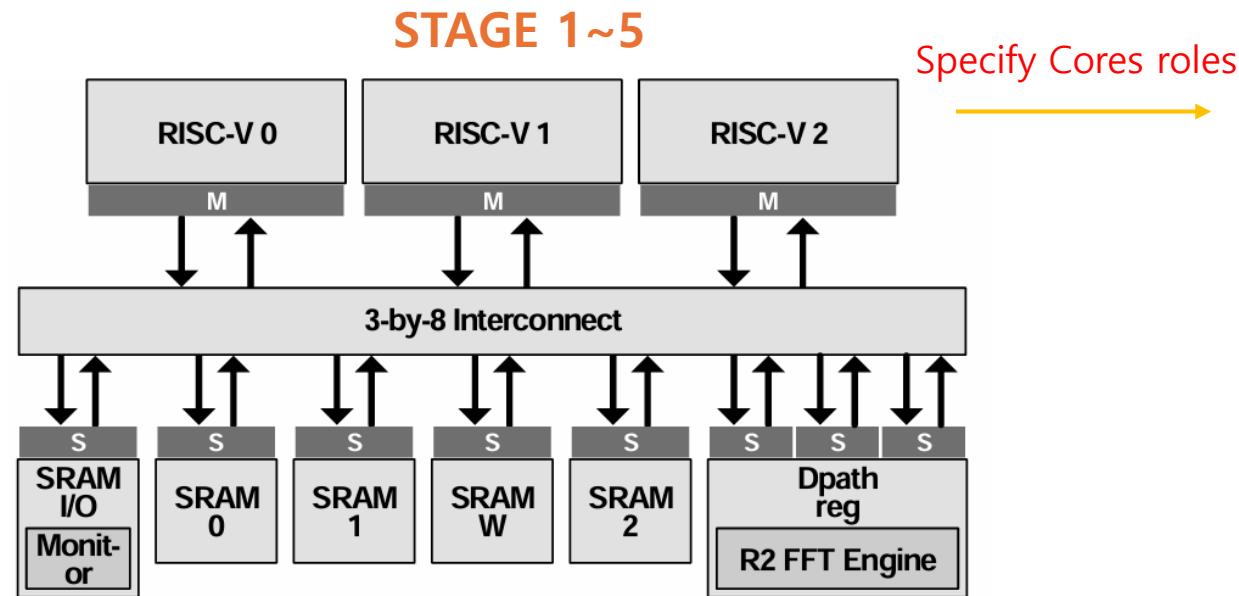
```
# Butterfly 0
lw t2, 0(s1)      # d0(SRAM_0->t2)
lw t3, 0(s2)      # d1(SRAM_1->t3)
sw t2, 0(s4)      # d0(t2->dpath0)
sw t3, 0(s5)      # d1(t3->dpath1)
addi x0, x0, 0x0000 # dummy
addi x0, x0, 0x0000 # dummy
```

Use 6 dummy codes before starts

Code Review (3/7)

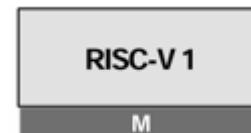
■ Implement POV-cont'd

- Pipelining with 3 cores (HOW?)
 - Distribute the roles among the cores



```
# Butterfly 0
lw t4, 0($3)      # W[0](SRAM_W->t4)
addi x0, x0, 1     # dummy
sw t4, 0($6)      # W[0](t4->dpath2)
lw t4, 64($3)     # W[16](SRAM_W->t4)_next twiddle factor
lw t5, 4($4)       # o_d0(dpath0->t5)
sw t5, 0($1)       # o_d0(t5->SRAM_0)
```

```
# Butterfly 1
sw t4, 0($6)      # W[16](t4->dpath2)
lw t4, 0($3)      # W[0](SRAM_W->t4)_next twiddle factor
lw t5, 4($4)       # o_d0(dpath0->t5)
sw t5, 0($2)       # o_d0(t5->SRAM_1)
```



```
# Butterfly operations (32 total)

# Butterfly 0
lw t6, 4($5)      # o_d1(dpath1->t6)
sw t6, 4($1)       # o_d1(t6->SRAM_1)
addi x0, x0, 1     # dummy
```



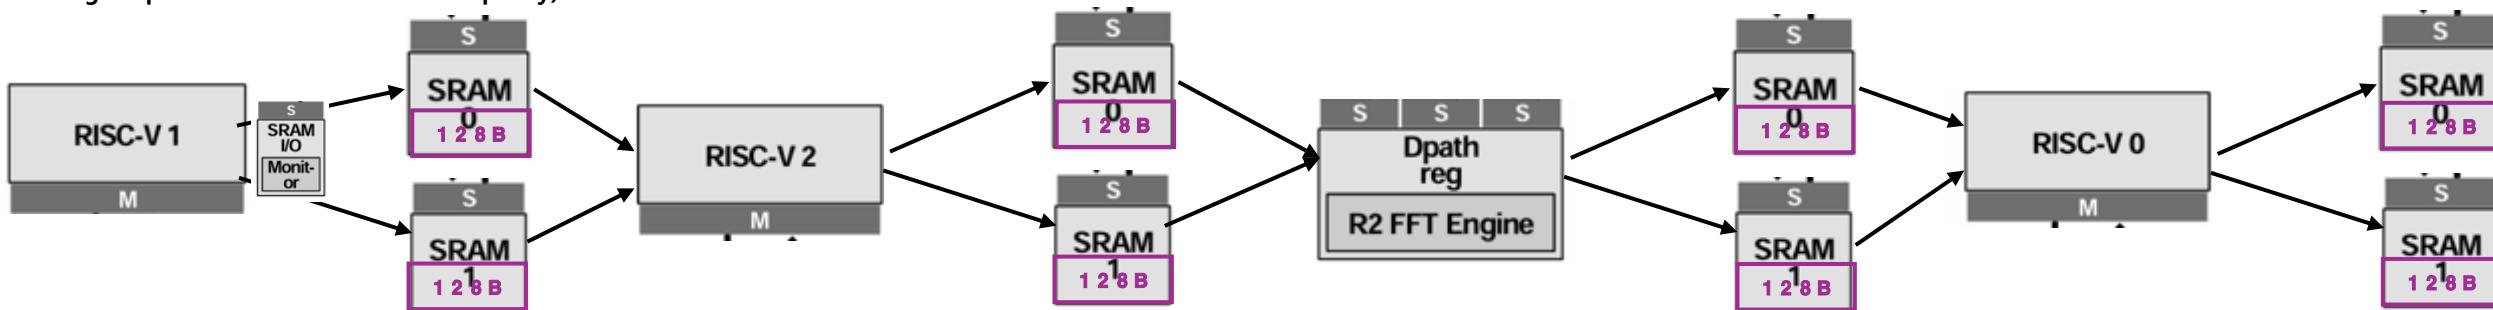
```
# Butterfly 0
lw t2, 0($1)       # d0(SRAM_0->t2)
lw t3, 4($1)       # d1(SRAM_0->t3)
sw t2, 0($4)       # d0(t2->dpath0)
sw t3, 0($5)       # d1(t3->dpath1)
```

Code Review (4/7)

■ Implement POV-cont'd

- Effective Memory Capacity
 - Memory is reused after use

(The "reg" expression was omitted for simplicity)



| | | |
|------|-------------|--------------------|
| lw | t0, 0(s0) | # d0(SRAM I/O->t0) |
| lw | t1, 128(s0) | # d1(SRAM I/O->t1) |
| sw | t0, 0(s1) | # d0(t0->SRAM_0) |
| sw | t1, 0(s2) | # d1(t1->SRAM_1) |
| addi | x0, x0, 1 | # dummy |

<Store Input to SRAM 0, 1>

| | | |
|------|----------------|------------------|
| lw | t2, 0(s1) | # d0(SRAM_0->t2) |
| lw | t3, 0(s2) | # d1(SRAM_1->t3) |
| sw | t2, 0(s4) | # d0(t2->dpath0) |
| sw | t3, 0(s5) | # d1(t3->dpath1) |
| addi | x0, x0, 0x0000 | # dummy |
| addi | x0, x0, 0x0000 | # dummy |

<Store dpath output to SRAM 0, 1>

| | | |
|----|-----------|--------------------|
| sw | t4, 0(s6) | # W[0](t4->dpath2) |
| lw | t5, 4(s4) | # o_d0(dpath0->t5) |
| lw | t6, 4(s5) | # o_d0(dpath1->t6) |
| sw | t5, 0(s1) | # o_d0(t5->SRAM_0) |
| sw | t6, 0(s2) | # o_d0(t6->SRAM_1) |

<Store dpath output to SRAM 0, 1>

Code Review (6/7)

▪ **Implement POV-cont'd**

- Synchronizing with dummy

- Dummy codes were inserted during IDLE state when a core had no tasks to perform

Synchronizing with other cores!!

```

lw      t4, 0($3)          # W[0](SRAM_W->t4)

addi   x0, x0, 1           # dummy

# Load twiddle factors for all stages (32 values)
# Butterfly 0-31: Load W[0] for stage 0 (all W[0])

# Butterfly 0
sw t4, 0($6)          # W[0](t4->dpath2)
lw t5, 4($4)          # o_d0(dpath0->t5)
lw t6, 4($5)          # o_d0(dpath1->t6)
sw t5, 0($1)          # o_d0(t5->SRAM_0)
sw t6, 0($2)          # o_d0(t6->SRAM_1)

```

Code Review (7/7)

■ Implement POV-cont'd

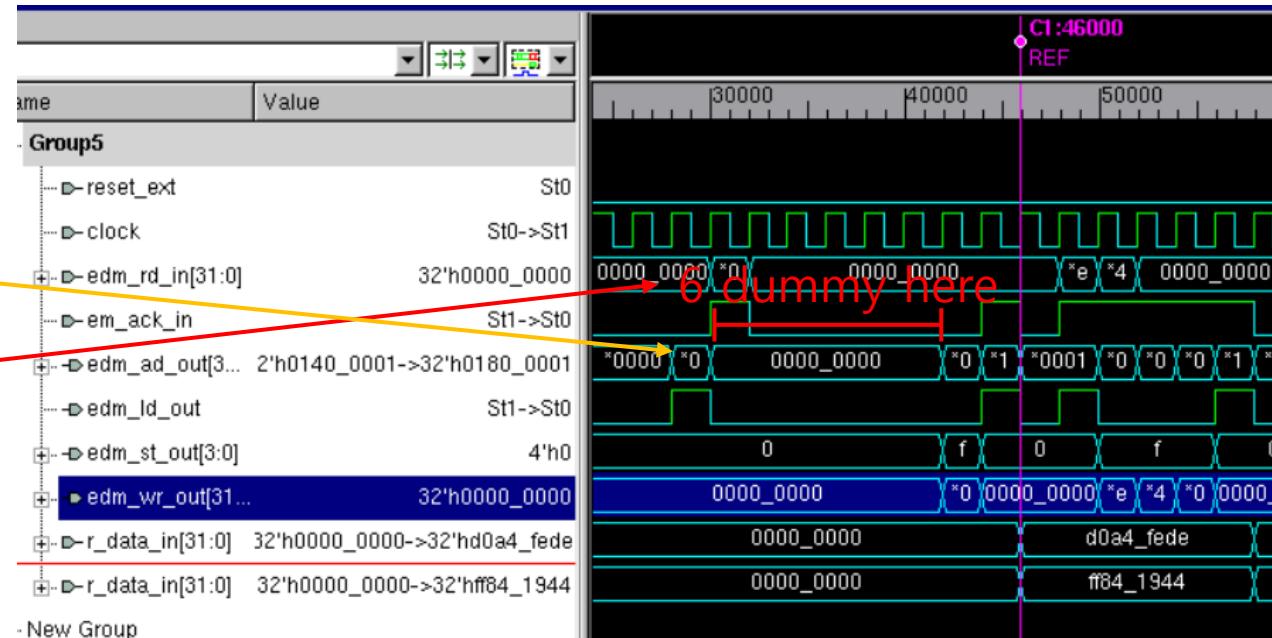
• Synchronizing with dummy

- Dummy codes were inserted during IDLE state when a core had no tasks to perform

```
lw      t4, 0(s3)      # W[0](SRAM_W->t4)
addi   x0, x0, 1        # dummy

# Load twiddle factors for all stages (32 values)
# Butterfly 0-31: Load W[0] for stage 0 (all W[0])

# Butterfly 0
sw t4, 0(s6)      # W[0](t4->dpath2)
lw t5, 4(s4)        # o_d0(dpath0->t5)
lw t6, 4(s5)        # o_d0(dpath1->t6)
sw t5, 0(s1)        # o_d0(t5->SRAM_0)
sw t6, 0(s2)        # o_d0(t6->SRAM_1)
```



How to Debug (1/7)

Using C & Excel for debugging with Waveforms

```
if (set == 0 && l == 1)
{
    printf("State of SRAM_0 after 0 stage of first set\n");
    for (j = 0; j < L_FFT; j++)
        printf("SRAM_0[%2d] = \t0x%8x\n", j, SRAM_0[j]);
}

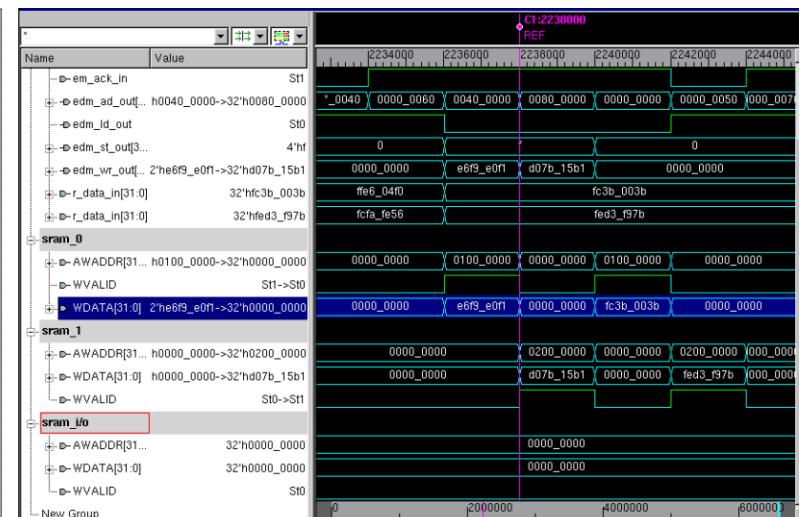
if (set == 0 && l == 2)
{
    printf("State of SRAM_0 after 1 stage of first set\n");
    for (j = 0; j < L_FFT; j++)
        printf("SRAM_0[%2d] = \t0x%8x\n", j, SRAM_0[j]);
}

if (set == 0 && l == 4)
{
    printf("State of SRAM_0 after 2 stage of first set\n");
    for (j = 0; j < L_FFT; j++)
        printf("SRAM_0[%2d] = \t0x%8x\n", j, SRAM_0[j]);
}

if (set == 0 && l == 8)
{
    printf("State of SRAM_0 after 3 stage of first set\n");
    for (j = 0; j < L_FFT; j++)
        printf("SRAM_0[%2d] = \t0x%8x\n", j, SRAM_0[j]);
}

if (set == 0 && l == 16)
{
    printf("State of SRAM_0 after 4 stage of first set\n");
    for (j = 0; j < L_FFT; j++)
        printf("SRAM_0[%2d] = \t0x%8x\n", j, SRAM_0[j]);
}
```

| stage0 | stage1 | stage2 | stage3 | stage4 | stage5 |
|-------------------------|--------------------------|--------------------------|--------------------------|--------------------------|--------------------------|
| SRAM_0[0] = 0xd004fe | SRAM_0[0] = 0xee7c4abc | SRAM_0[0] = 0xed3fb607 | SRAM_0[0] = 0xf515fe6e | SRAM_0[0] = 0xfacaf646 | SRAM_0[0] = 0xfb0ff96 |
| SRAM_0[1] = 0xff841944 | SRAM_0[1] = 0xc54ea9a | SRAM_0[1] = 0xe2280a22 | SRAM_0[1] = 0x9f9d103d | SRAM_0[1] = 0x44a1692 | SRAM_0[1] = 0xdf6706e0 |
| SRAM_0[2] = 0x5d5e1278 | SRAM_0[2] = 0x5f5e0f3 | SRAM_0[2] = 0x0f68fc7 | SRAM_0[2] = 0x7f70043f | SRAM_0[2] = 0x3cf9de2 | SRAM_0[2] = 0x35cfb3c |
| SRAM_0[3] = 0xd5e1278 | SRAM_0[3] = 0x6f86051 | SRAM_0[3] = 0xd0406f7 | SRAM_0[3] = 0xfe7809ad | SRAM_0[3] = 0x8200af | SRAM_0[3] = 0xdfdbf023 |
| SRAM_0[4] = 0xfc2ffff | SRAM_0[4] = 0xebf90152 | SRAM_0[4] = 0x1419b5 | SRAM_0[4] = 0x768d44 | SRAM_0[4] = 0x4adff93 | SRAM_0[4] = 0x1330140 |
| SRAM_0[5] = 0xd8f83fec | SRAM_0[5] = 0xe3dfe60a | SRAM_0[5] = 0x25fdb5 | SRAM_0[5] = 0x96eff9d | SRAM_0[5] = 0x8a6ff70 | SRAM_0[5] = 0x4e20010 |
| SRAM_0[6] = 0xdbcb402a | SRAM_0[6] = 0x1035fe8 | SRAM_0[6] = 0x1f1c0042b | SRAM_0[6] = 0xf52e0a27 | SRAM_0[6] = 0x3b7054d | SRAM_0[6] = 0x2e70157 |
| SRAM_0[7] = 0x296e62a | SRAM_0[7] = 0xbfbac002a | SRAM_0[7] = 0x9f620459 | SRAM_0[7] = 0x2c107a1 | SRAM_0[7] = 0x0fa9305e9 | SRAM_0[7] = 0x0fad4001e |
| SRAM_0[8] = 0xf2b20ced | SRAM_0[8] = 0xc0c0f69 | SRAM_0[8] = 0xfc001d7 | SRAM_0[8] = 0x825f68 | SRAM_0[8] = 0x0fd802b4 | SRAM_0[8] = 0x2f04411 |
| SRAM_0[9] = 0x5a022a3 | SRAM_0[9] = 0x5e2376 | SRAM_0[9] = 0x332209d | SRAM_0[9] = 0xf5349aa | SRAM_0[9] = 0xdaf000d8 | SRAM_0[9] = 0x0fc0f0357 |
| SRAM_0[10] = 0x46711e6 | SRAM_0[10] = 0x7af6d83 | SRAM_0[10] = 0x7f7007c8 | SRAM_0[10] = 0x8f78ff87 | SRAM_0[10] = 0x4fffa1 | SRAM_0[10] = 0x3ecfaee |
| SRAM_0[11] = 0xbdb71b1c | SRAM_0[11] = 0x542f2c2 | SRAM_0[11] = 0xe45040aa | SRAM_0[11] = 0xe8fcdf49 | SRAM_0[11] = 0x0feb9ff24 | SRAM_0[11] = 0x0fe73ff9e |
| SRAM_0[12] = 0xebc7eb91 | SRAM_0[12] = 0xffffd446 | SRAM_0[12] = 0x1f1c0d91 | SRAM_0[12] = 0x9f9d8f66 | SRAM_0[12] = 0x1fe16fe80 | SRAM_0[12] = 0xfcdd1ff34 |
| SRAM_0[13] = 0x55d14cd | SRAM_0[13] = 0x3f36194e | SRAM_0[13] = 0x0fd2c02d9 | SRAM_0[13] = 0x5b6fe17 | SRAM_0[13] = 0x3f0470 | SRAM_0[13] = 0x0fe5f02d5 |
| SRAM_0[14] = 0xe2fcfb | SRAM_0[14] = 0x0edf2f74a | SRAM_0[14] = 0x2ef5b | SRAM_0[14] = 0xfc920303 | SRAM_0[14] = 0x0fbabfd62 | SRAM_0[14] = 0x0ff5f050c |
| SRAM_0[15] = 0xe230f911 | SRAM_0[15] = 0xf626fb7e | SRAM_0[15] = 0x0fb0dfa2 | SRAM_0[15] = 0x0f6c1fcb7 | SRAM_0[15] = 0x0fd9ff70 | SRAM_0[15] = 0x0f8a4fc40 |
| SRAM_0[16] = 0x1b101afa | SRAM_0[16] = 0xb5f1fd8b | SRAM_0[16] = 0x0f7a7ff47 | SRAM_0[16] = 0x81fc5a | SRAM_0[16] = 0x0fa4a010a | SRAM_0[16] = 0x0fcfe0fe |
| SRAM_0[17] = 0xffffd27a | SRAM_0[17] = 0xb33fc45 | SRAM_0[17] = 0x0f63003b3 | SRAM_0[17] = 0x0ff3df55 | SRAM_0[17] = 0x2c30ba8 | SRAM_0[17] = 0x63a0516 |
| SRAM_0[18] = 0xbcb2e023 | SRAM_0[18] = 0x1fb1f67 | SRAM_0[18] = 0x1120f6c0 | SRAM_0[18] = 0x11b0fe25 | SRAM_0[18] = 0x0f427065c | SRAM_0[18] = 0x0f58104f7 |
| SRAM_0[19] = 0xb8b0697 | SRAM_0[19] = 0x49df595 | SRAM_0[19] = 0x6849f08 | SRAM_0[19] = 0x6bcfa82 | SRAM_0[19] = 0x0fd608fd | SRAM_0[19] = 0x407b9 |
| SRAM_0[20] = 0x3b03a6 | SRAM_0[20] = 0x5f5e0105 | SRAM_0[20] = 0x3a9fe43 | SRAM_0[20] = 0x5b9d07 | SRAM_0[20] = 0x2b2bd1b9 | SRAM_0[20] = 0x0ff960103 |
| SRAM_0[21] = 0xea010566 | SRAM_0[21] = 0xe18fc4 | SRAM_0[21] = 0x502f931 | SRAM_0[21] = 0x6aa050e | SRAM_0[21] = 0xc3012c | SRAM_0[21] = 0x2a6ff3a |
| SRAM_0[22] = 0x418f64 | SRAM_0[22] = 0xe602a1 | SRAM_0[22] = 0x6f816a6 | SRAM_0[22] = 0x691109 | SRAM_0[22] = 0x0f17704d9 | SRAM_0[22] = 0x0f76105d6 |
| SRAM_0[23] = 0xb7f190 | SRAM_0[23] = 0xfc808072 | SRAM_0[23] = 0x1fe18fc8 | SRAM_0[23] = 0x9f93b37b | SRAM_0[23] = 0x82d01b7 | SRAM_0[23] = 0x205015e |
| SRAM_0[24] = 0x1ac7fd9c | SRAM_0[24] = 0xd930138 | SRAM_0[24] = 0x95df96e | SRAM_0[24] = 0x0f72502e | SRAM_0[24] = 0x0fa9cf9de | SRAM_0[24] = 0x0fcff7f6a |
| SRAM_0[25] = 0x1029141e | SRAM_0[25] = 0x6fc5ff6 | SRAM_0[25] = 0x994fe86 | SRAM_0[25] = 0xf6f3045d | SRAM_0[25] = 0x0f79348d2 | SRAM_0[25] = 0x0fa2efae8 |
| SRAM_0[26] = 0x6104d5 | SRAM_0[26] = 0x3d3fc64 | SRAM_0[26] = 0x14430579 | SRAM_0[26] = 0x0ff7e089b | SRAM_0[26] = 0x0f829fd5 | SRAM_0[26] = 0x0fc650537 |
| SRAM_0[27] = 0x14a7dd63 | SRAM_0[27] = 0x19631462 | SRAM_0[27] = 0x66004e4 | SRAM_0[27] = 0xffc8fe86 | SRAM_0[27] = 0x0ffd3fe4 | SRAM_0[27] = 0x0fae7fe0a |



How to Debug (2/7)

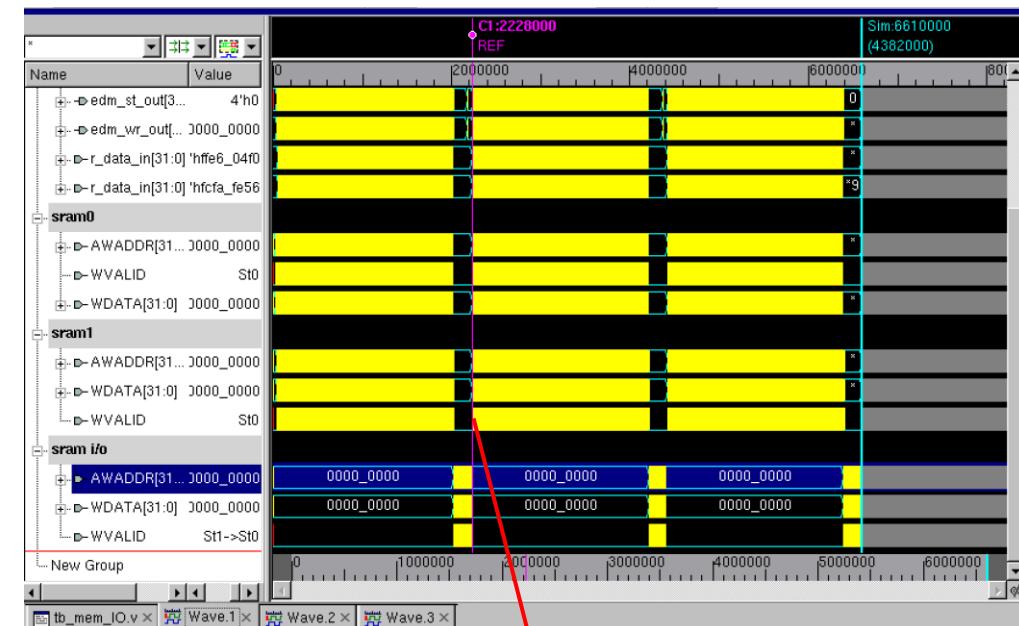
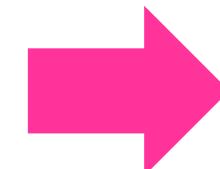
■ Debugging Example (Timing analysis)

- See the output result (console)

Output 0~63(set0)correct, 64~ wrong..

Output count has reached the total number of samples.
Validating bit-accurate output.
- Mismatch @ output[0x0040]: expected 0xf551fc78, got 0xfd23fa07
- Mismatch @ output[0x0041]: expected 0x0184f8db, got 0xfe450daf
- Mismatch @ output[0x0042]: expected 0xfcce5fe76, got 0x00f7fbef
- Mismatch @ output[0x0043]: expected 0xfe16f4f1, got 0x0169119d
- Mismatch @ output[0x0044]: expected 0x03a4fd74, got 0xfe7a00ae
- Mismatch @ output[0x0045]: expected 0x0561fd93, got 0xf522004b

See the set1_stage 0

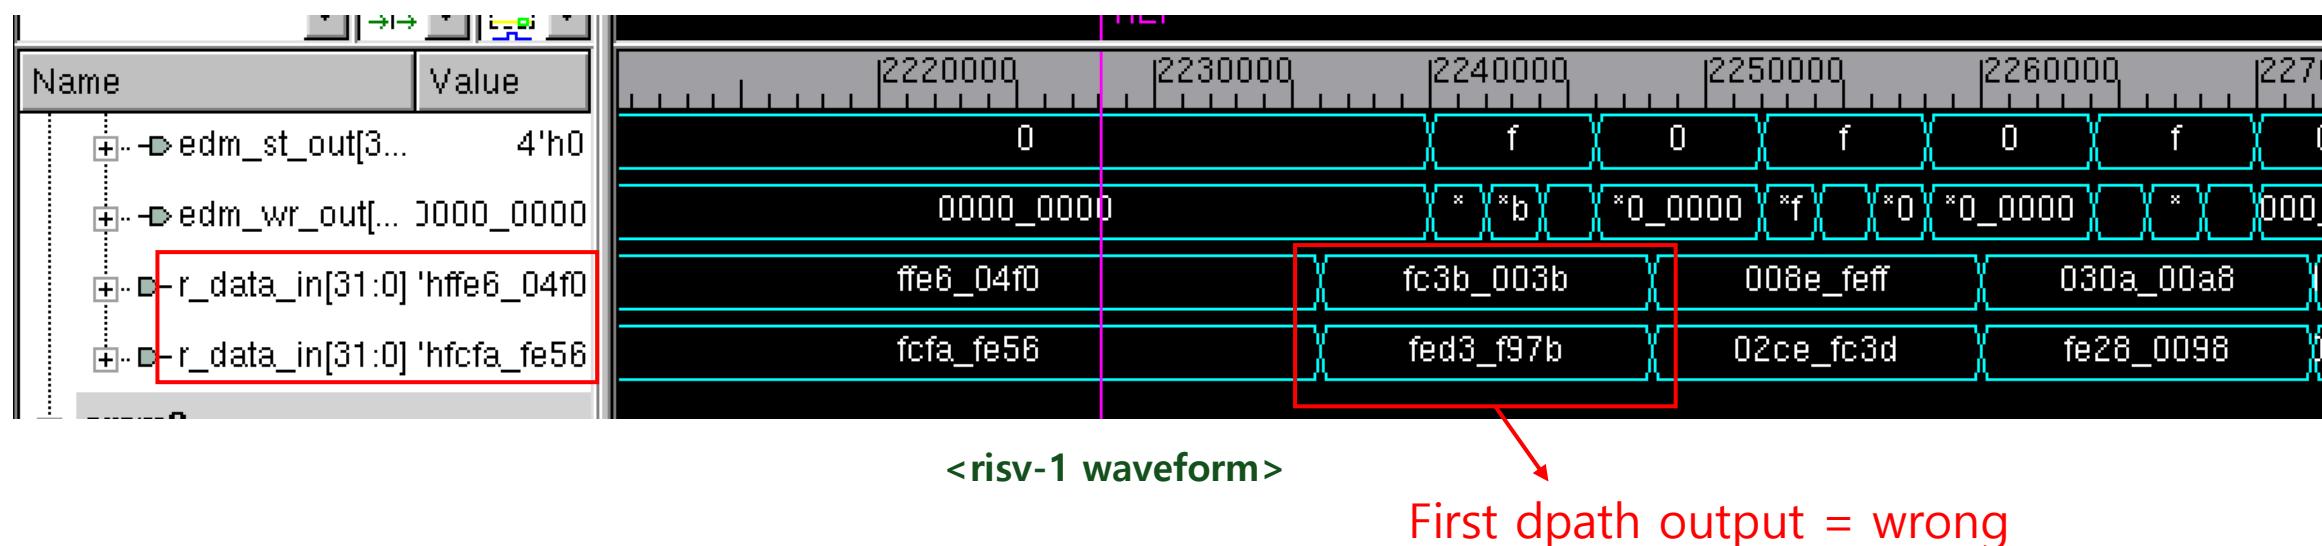


Set1_stage 0

How to Debug (4/7)

■ Debugging Example (Timing analysis)-cont'd

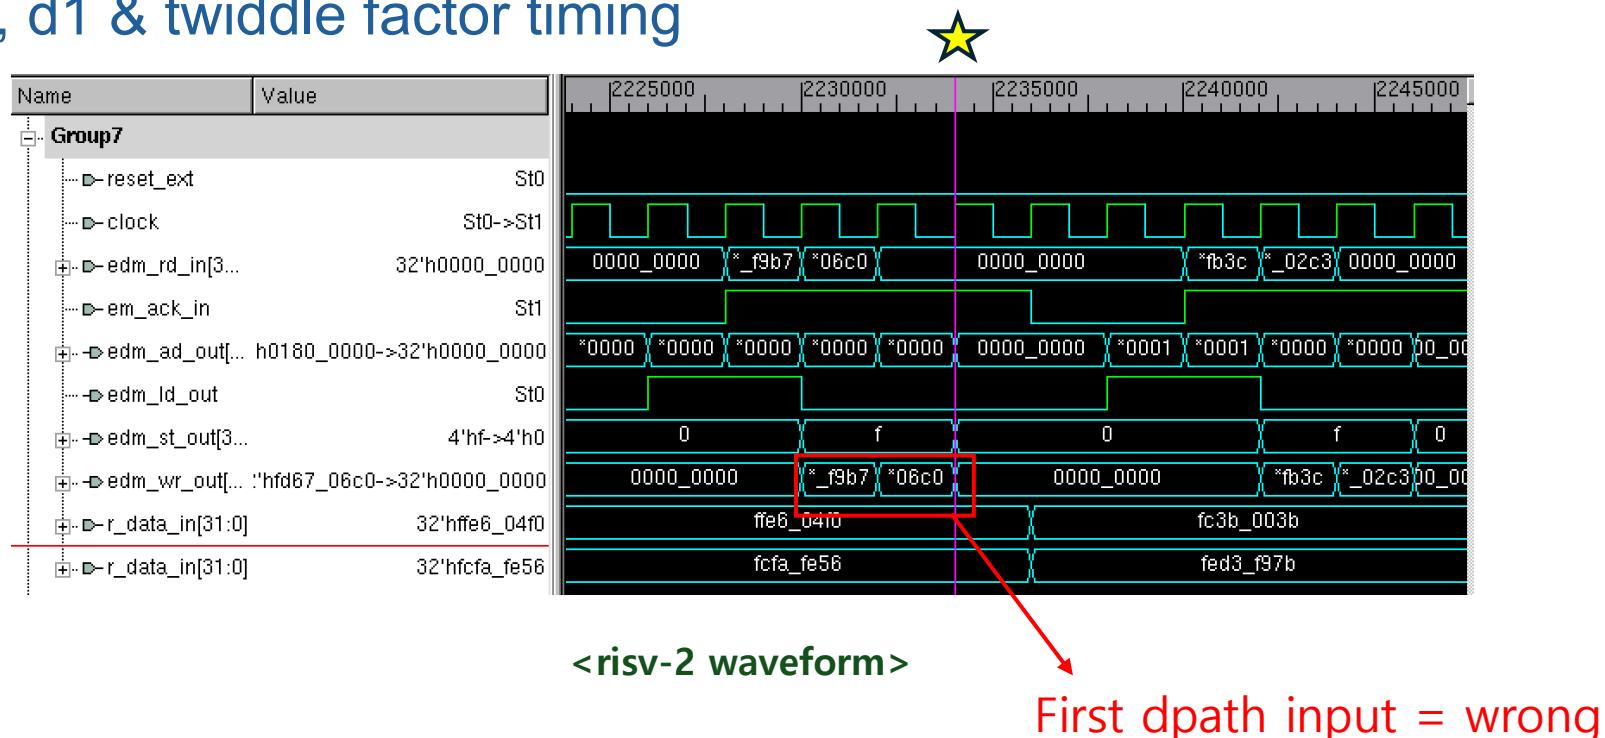
- See the d0, d1 & twiddle factor timing



How to Debug (5/7)

■ Debugging Example (Timing analysis)-cont'd

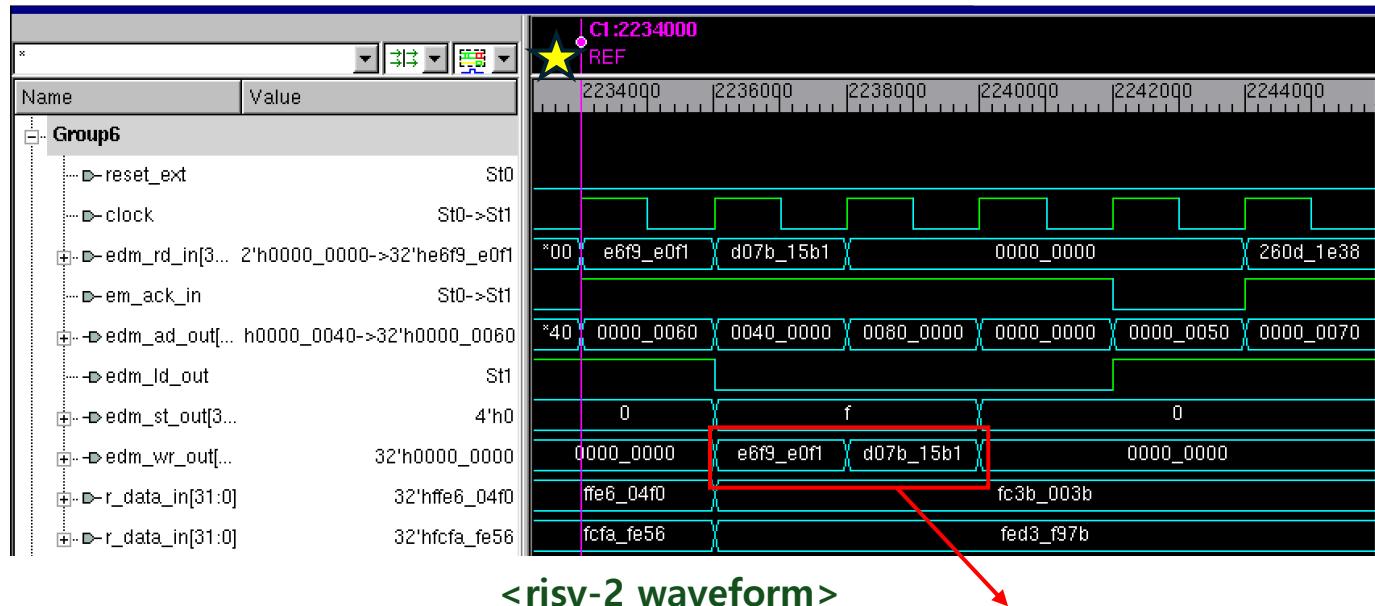
- See the d0, d1 & twiddle factor timing



How to Debug (6/7)

■ Debugging Example (Timing analysis)-cont'd

- See the input timing



<risv-2 waveform>

Right input is to late
=> Compute core (risv-1) need stall

How to Debug (7/7)

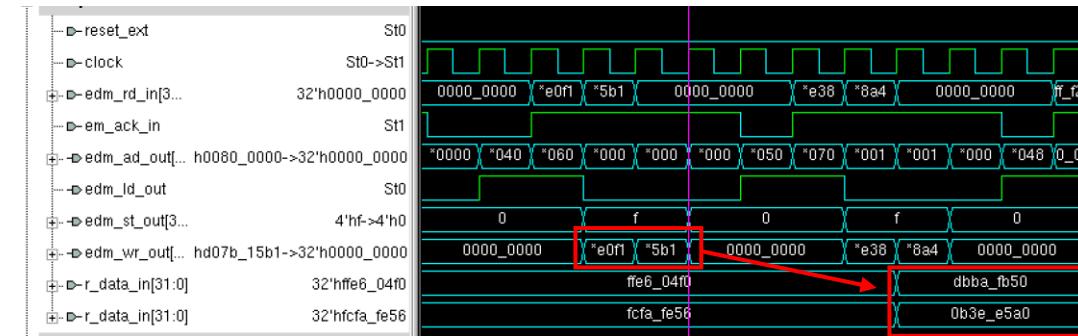
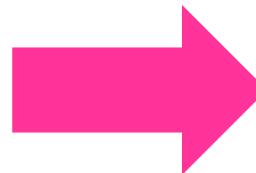
■ Debugging Example (Timing analysis)-cont'd

- Add 6 dummy(nop) to data_m2.s and check

```
#butterfly 31
lw s11, 124(s1) #sram_0[31]->s11
sw s11, 248(s7) #s11->sram i/o
addi x0, x0, 0x0000
# dummy
addi x0, x0, 0x0000
```

```
# ===== set0 End =====
```

<data_m2.s>



<risv-1 waveform>

Right timing! => correct result!!

```
Output count has reached the total number of samples.
Validating bit-accurate output...
- Total mismatches: 0 / 192
Checking SRAM occupancy...
- SRAM 2: occupancy = 32
- SRAM 3: occupancy = 32
Test passed. Ending simulation.
$finish called from file "./tb/tb_mem_ID.v", line 132.
$finish at simulation time 6634000
Simulation complete, time is 6634000 ps.
```

<console>

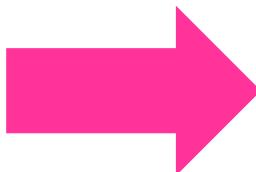
Pipelining With Memory Flag

- Code for Flagging
- Execution time & Memory Capacity

```
# Clear sync flags initially
sw    x0,    0($2)      # Clear ready flag
sw    x0,    4($2)      # Clear done flag

main_loop:
bge   s5,    s4,    end_infinite

# Wait for Core 2 to be ready (check if previous output is done)
beq   s5,    x0,    skip_wait  # Skip wait for first set
wait_core2:
lw    t0,    4($2)      # Check done flag from Core 2
beq   t0,    x0,    wait_core2
sw    x0,    4($2)      # Clear done flag
```



```
Output count has reached the total number of samples.
Validating bit-accurate output...
- Total mismatches: 0 / 192
Checking SRAM occupancy...
- SRAM 2: occupancy = 64
- SRAM 3: occupancy = 3
Test passed. Ending simulation.
$finish called from file "../../tb/tb_mem_IO.v", line 132.
$finish at simulation time 60022000
Simulation complete, time is 60022000 ps.
```

Performance

- Simulation Time : 60,022,000[ps]

Memory Capacity

- Used Memory: SRAM 0: 64(128B)
- Used Memory: SRAM 1: 3(128B)

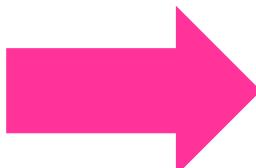
SRAM 1 for Flags

■ Parallel Operations /w TP_1 code

- Reuse TP_1 code & dummy for Sync

```
# reg에 Twiddle Factor을 담아두고, 재사용하려한다.  
lw      a7, 0($3)          # a7 = twiddle[0]  
lw      a6, 64($3)         # a6 = twiddle[16]  
  
# i = 0  
# n = 0  
# stage[0][0]  
  
# stage[0][0]의 연산 결과를 stage[0][1]로 전달하는데  
# SRAM 0의 sw/lw 과정을 제하기위해 가용한 reg를 최대한 사용하였다.  
  
# Input  
# 아래는 DIT 형식에 맞추어 SRAM I/O으로부터의 Input을 받는 과정이다.  
# e.g. x[0], x[32], x[16]...순서...  
lw      t0, 0($0)  
lw      t1, 128($0)  
lw      t2, 64($0)  
lw      t3, 192($0)  
lw      t4, 32($0)  
lw      t5, 160($0)  
lw      t6, 96($0)  
lw      a0, 224($0)
```

- Execution time & Memory Capacity



```
Output count has reached the total number of samples.  
Validating bit-accurate output...  
- Total mismatches: 0 / 192  
Checking SRAM occupancy...  
- SRAM 2: occupancy = 64  
- SRAM 3: occupancy = 64  
- SRAM 5: occupancy = 64  
Test passed. Ending simulation.  
$finish called from file "../../../../tb/tb\_mem\_I0.v", line 132.  
$finish at simulation time 6044000  
Simulation complete, time is 6044000 ps.
```

Performance X 10!!

■ Performance

- Simulation Time : 6,044,000[ps]

■ Memory Capacity

- Used Memory: SRAM 0: 64(128B)
- Used Memory: SRAM 1: 64(128B)
- Used Memory: SRAM 2: 64(128B)

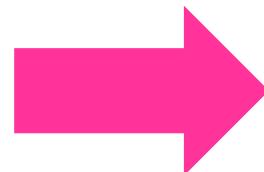
■ Pipelining Cores

- code

```
stage0:  
    # Initial synchronization - wait for m1 to start loading  
  
    lw      t4, 0($3)      # W[0](SRAM_W->t4)  
  
    addi   x0, x0, 1       # dummy  
    addi   x0, x0, 1       # dummy  
  
    # Load twiddle factors for all stages (32 values)  
    # Butterfly 0-31: Load W[0] for stage 0 (all W[0])  
  
    # Butterfly 0  
    sw t4, 0($6)          # W[0](t4->dpah2)  
    lw t5, 4($4)          # o_d0(dpah0->t5)  
    lw t6, 4($5)          # o_d0(dpah1->t6)  
    sw t5, 0($1)          # o_d0(t5->SRAM_0)  
    sw t6, 0($2)          # o_d0(t6->SRAM_1)
```

Memory X 3!!

- Execution time & Memory Capacity



```
Output count has reached the total number of samples.  
Validating bit-accurate output...  
    - Total mismatches: 0 / 192  
Checking SRAM occupancy...  
    - SRAM 2: occupancy = 32  
    - SRAM 3: occupancy = 32  
Test passed. Ending simulation.  
$finish called from file "../../tb/tb\_mem\_IO.v", line 132.  
$finish at simulation time 6804000  
Simulation complete, time is 6804000 ps.
```

■ Performance

- Simulation Time : 6,804,000[ps]

■ Memory Capacity

- Used Memory: SRAM 0: 32(128B)
- Used Memory: SRAM 1: 32(128B)

Ablation Study(4/4)

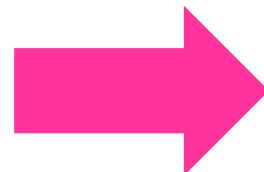
■ Pipeling Cores (Optimized)-final version

- code

```
stage0:  
    # Initial synchronization - wait for m1 to start loading  
  
    lw      t4, 0($3)      # W[0](SRAM_W->t4)  
  
    addi   x0, x0, 1       # dummy  
    addi   x0, x0, 1       # dummy  
  
    # Load twiddle factors for all stages (32 values)  
    # Butterfly 0-31: Load W[0] for stage 0 (all W[0])  
  
    # Butterfly 0  
    sw t4, 0($6)          # W[0](t4->dpah2)  
    lw t5, 4($4)          # o_d0(dpah0->t5)  
    lw t6, 4($5)          # o_d0(dpah1->t6)  
    sw t5, 0($1)          # o_d0(t5->SRAM_0)  
    sw t6, 0($2)          # o_d0(t6->SRAM_1)
```

Performance X 1.1 !!

- Execution time & Memory Capacity



```
Output count has reached the total number of samples.  
Validating bit-accurate output...  
- Total mismatches: 0 / 192  
Checking SRAM occupancy...  
- SRAM 2: occupancy = 32  
- SRAM 3: occupancy = 32  
Test passed. Ending simulation.  
$finish called from file "../../../../tb/tb\_mem\_IO.v", line 132.  
$finish at simulation time 6634000  
Simulation complete, time is 6634000 ps.
```

■ Performance

- Simulation Time : 6,634,000[ps]

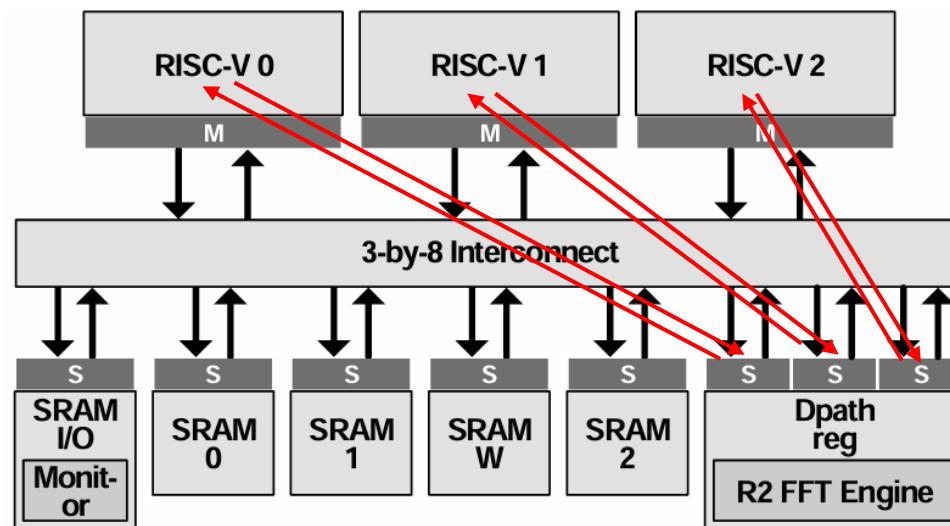
■ Memory Capacity

- Used Memory: SRAM 0: 32(128B)
- Used Memory: SRAM 1: 32(128B)

Roofline Model(1/4)

■ System Analysis

- Roofline analysis



Y axis (perf)
X axis (IO)

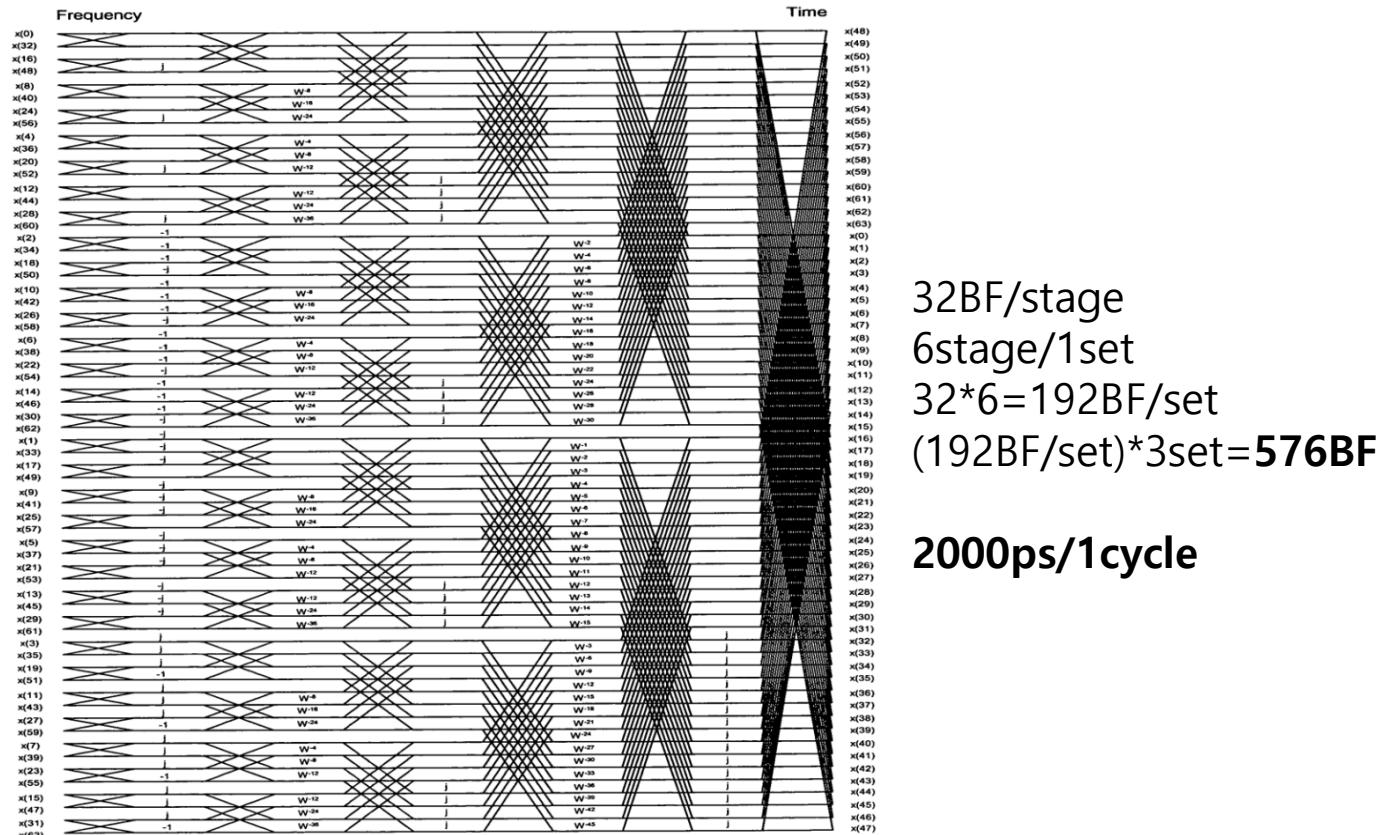
Performance limit
BW_m (Slope)

: BF/cycle
: BF/BYTE

: 1BF/cycle
: 12Byte/cycle (3core)

Roofline Model(2/4)

■ Compute Performance (BF/Cycle)



- Pipelining with memory flag
 $(60022000 \text{ ps}) * (1\text{cycle}/2000\text{ps}) = 30,011\text{cycle}$
576BF/30,011cycle = **0.019 [BF/cycle]**
- Parallel Operations /w TP_1 code
 $(6044000 \text{ ps}) * (1\text{cycle}/2000\text{ps}) = 3,022\text{cycle}$
576BF/3,022cycle = **0.191 [BF/cycle]**
- Pipelining Cores
 $(6804000 \text{ ps}) * (1\text{cycle}/2000\text{ps}) = 3,402\text{cycle}$
576BF/3,402cycle = **0.169 [BF/cycle]**
- Final Ver.
 $(6634000 \text{ ps}) * (1\text{cycle}/2000\text{ps}) = 3,317\text{cycle}$
576BF/3,317cycle = **0.174 [BF/cycle]**

■ Compute IO (BF/Byte)-cont'd

$1\text{BF} \Rightarrow 3\text{LW} + 3\text{SW}$ (input), $2\text{LW} + 2\text{SW}$ (output)

$$\text{IO} = 1\text{BF}/(10*4)\text{BYTE} = 0.025(\text{BF/BYTE})$$

- Pipelining with memory flag

-Total

Input load : $192*4$ Byte
Flag check : $(6*3*4)*576$ Byte (plag poling)
Compute : $10*4*576$ Byte
Output store : $192*4$ Byte

$$\begin{aligned}\text{IO} &= 576\text{BF} / (192*4 + 6*3*4*576 + \\ &\quad 10*4*576 + 192*4) \\ &= 576\text{BF} / 66048\text{Byte} \\ &= 0.0087(\text{BF/BYTE})\end{aligned}$$

- Parallel Operations /w TP_1 code

-Total

Input load : $192*4$ Byte
Compute(d_1, d_2) : $8*4*576$ Byte
Compute(W) : $(1+1+2+4+8+24+576)*4$ Byte
Output store : $192*4$ Byte

$$\begin{aligned}\text{IO} &= 576\text{BF} / 22432\text{Byte} \\ &= 0.0257(\text{BF/BYTE})\end{aligned}$$

Roofline Model(4/4)

■ Compute IO (BF/Byte)-cont'd

- Pipelining Cores & Final ver.

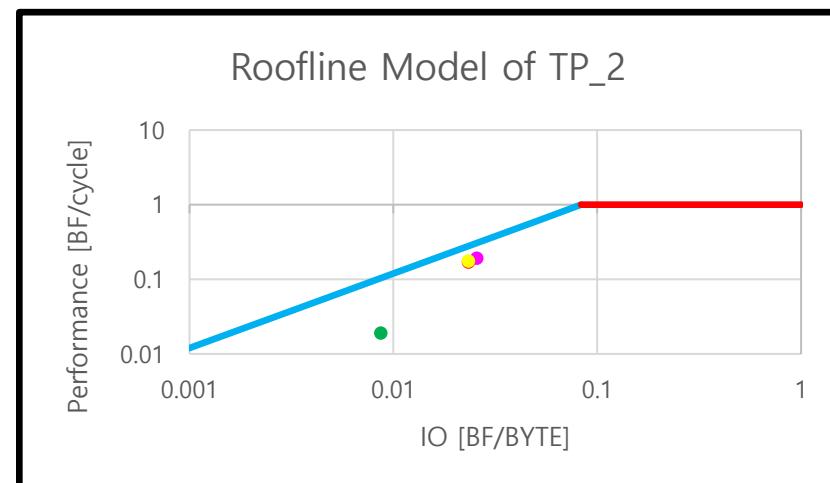
-Total

Input load : 192×4 Byte
Compute : $10 \times 4 \times 576$ Byte
Output store : 192×4 Byte

$$\text{IO} = 576\text{BF} / 24576 \text{ Byte}$$
$$= 0.0234(\text{BF/BYTE})$$

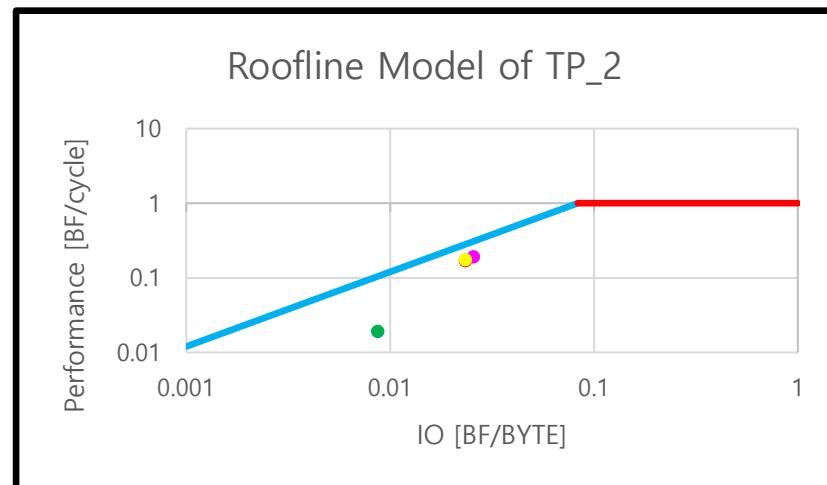
■ Draw roofline model /w Excel

| A | B | C | D | E | F | G |
|----------------------------------|---------|-----------------------|-------------|--------------|------------|------------|
| | BF/byte | Performance(BF/cycle) | Memory line | Compute line | X_roofline | Y_roofline |
| 2 Pipelining with memory flag | | 0.0087 | 0.019 | 0.1044 | 1 | 0.001 |
| 3 Parallel Operations /w TP_1 cc | | 0.0257 | 0.191 | 0.3084 | 1 | 0.005 |
| 4 Pipelining Cores | | 0.0234 | 0.169 | 0.2808 | 1 | 0.01 |
| 5 Final ver. | | 0.0234 | 0.174 | 0.2808 | 1 | 0.05 |
| | | | | | 0.0833 | 1 |
| | | | | | 0.1 | 1 |
| | | | | | 0.5 | 1 |
| | | | | | 1 | 1 |

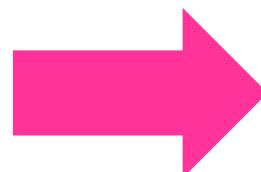


Our system is Memory Bound !!

How to improve



Our system is Memory Bound !!



- Reuse twiddle factor
- Using register
- Etc..

How to improve



: core 0 is wasted..

If core 0 do LW & SW Next set input data?



core 1, 2 are good!!



■ Compare with TP_1

- 3 cores vs Single core

| | Performance[ps] | Memory Capacity[B/B] |
|--------------------|-----------------|----------------------|
| Single Core (TP_1) | 14,190,000[ps] | 256B/256B |
| 3 Cores (TP_2) | 6,634,000[ps] | 256B/768B |